

Step 1: Create the **employee** and **office** tables

Before applying the joins, we'll first create and insert records into the **employee** and **office** tables:

-- Create employee table

```
CREATE TABLE employee (  
  employee_id INT PRIMARY KEY, -- Unique identifier for each employee  
  first_name VARCHAR(50),      -- Employee's first name  
  last_name VARCHAR(50),       -- Employee's last name  
  office_id INT                -- The office where the employee works (linked to office table)  
);
```

-- Create office table

```
CREATE TABLE office (  
  office_id INT PRIMARY KEY, -- Unique identifier for each office  
  city VARCHAR(50),          -- The city where the office is located  
  country VARCHAR(50)        -- The country where the office is located  
);
```

-- Insert records into employee table

```
INSERT INTO employee (employee_id, first_name, last_name, office_id) VALUES  
(1, 'John', 'Doe', 101),      -- Employee John Doe works in office 101  
(2, 'Jane', 'Smith', 102),    -- Employee Jane Smith works in office 102  
(3, 'Michael', 'Johnson', 101), -- Employee Michael Johnson works in office 101  
(4, 'Emily', 'Davis', 103),    -- Employee Emily Davis works in office 103  
(5, 'Robert', 'Brown', NULL);  -- Robert Brown has no office assigned (NULL office_id)
```

-- Insert records into office table

```
INSERT INTO office (office_id, city, country) VALUES  
(101, 'New York', 'USA'),      -- Office 101 is in New York, USA  
(102, 'London', 'UK'),         -- Office 102 is in London, UK  
(103, 'Berlin', 'Germany');    -- Office 103 is in Berlin, Germany
```

1. INNER JOIN

Returns records that have matching values in both tables.

```
SELECT e.first_name, e.last_name, o.city, o.country  
FROM employee e  
INNER JOIN office o ON e.office_id = o.office_id;
```

Explanation:

- `SELECT e.first_name, e.last_name, o.city, o.country:`
 - Selects four columns:
 - `e.first_name`: First name from the `employee` table.
 - `e.last_name`: Last name from the `employee` table.
 - `o.city`: City from the `office` table.
 - `o.country`: Country from the `office` table.
- `FROM employee e:`
 - Specifies that the `employee` table is being used as the main table (aliased as `e` for shorter reference).
- `INNER JOIN office o:`
 - Joins the `employee` table (`e`) with the `office` table (`o`).
 - `INNER JOIN` ensures that only rows with matching `office_id` values in both tables are included in the result.
- `ON e.office_id = o.office_id:`
 - Defines the condition for the join, where `office_id` in the `employee` table must match `office_id` in the `office` table.
 - If an `employee` has an `office_id` that exists in the `office` table, their information is included.

Result:

- Employees with an assigned office (John, Jane, Michael, Emily) will be shown with their office locations. Employees without an office (Robert) are excluded.

1. **LIKE** Condition:

- The **LIKE** condition is used in string pattern matching. It can be used in joins to match columns with similar patterns.

Example:

Suppose you have a column `first_name` in the `employee` table, and you want to join only those rows where the first name starts with 'J'.

```
SELECT e.first_name, e.last_name, o.city
```

```
FROM employee e
```

```
INNER JOIN office o ON e.office_id = o.office_id
```

```
WHERE e.first_name LIKE 'J%';
```

Explanation:

- **LIKE 'J%'**: This condition matches employees whose first name starts with the letter 'J'.
- The join returns only employees whose first name starts with 'J', along with their office details.

2. **IN** Condition:

- The **IN** condition is used to check if a column value exists within a specified list of values. It can also be used within joins.

Example:

If you want to select employees who work in specific offices (say, offices with **office_id** 101 and 103):

```
SELECT e.first_name, e.last_name, o.city
FROM employee e
INNER JOIN office o ON e.office_id = o.office_id
WHERE e.office_id IN (101, 103);
```

Explanation:

- **WHERE e.office_id IN (101, 103)**: This condition filters only those employees whose **office_id** is either 101 or 103.
- It combines with the **INNER JOIN** to return matching employees who work in New York (101) and Berlin (103).

3. **IS NULL / IS NOT NULL** Condition:

- Used to filter records where certain columns are either **NULL** or not **NULL**.

Example:

Suppose you want to join all employees and return only those who have not been assigned an office (**office_id** is **NULL**).\

```
SELECT e.first_name, e.last_name  
  
FROM employee e  
  
LEFT JOIN office o ON e.office_id = o.office_id  
  
WHERE e.office_id IS NULL;
```

Explanation:

- **WHERE e.office_id IS NULL:** This condition filters only those employees who don't have an assigned **office_id** (i.e., the **office_id** is **NULL**).
- The **LEFT JOIN** ensures that all employees are included, and this condition filters for those with no office.

2. LEFT JOIN (or LEFT OUTER JOIN):

Returns all records from the left table (**employee**), and the matched records from the right table (**office**). If there is no match, **NULL** is returned for the right table columns.

```
SELECT e.first_name, e.last_name, o.city, o.country  
FROM employee e  
LEFT JOIN office o ON e.office_id = o.office_id;
```

Explanation:

- **SELECT e.first_name, e.last_name, o.city, o.country:**
 - Similar to the **INNER JOIN**, selects first name, last name, city, and country.
- **FROM employee e:**
 - The **employee** table is again the main table (aliased as **e**).
- **LEFT JOIN office o:**
 - Joins the **employee** table with the **office** table.
 - **LEFT JOIN** includes all rows from the **employee** table, regardless of whether there is a matching **office_id** in the **office** table.
- **ON e.office_id = o.office_id:**
 - The condition for the join. If the **employee** has no **office_id** (e.g., Robert Brown has **NULL**), **NULL** will be returned for the **city** and **country** from the **office** table.

Result:

- All employees (including Robert who has no office) will be shown. For employees without an office, `city` and `country` will be `NULL`.

2. Range Condition (**BETWEEN**):

- Used to filter records by a range of values.

```
SELECT e.first_name, e.last_name, o.city, o.country
FROM employee e
LEFT JOIN office o ON e.office_id = o.office_id
WHERE e.employee_id BETWEEN 2 AND 4;
```

3. RIGHT JOIN (or RIGHT OUTER JOIN)

Returns all records from the right table (`office`), and the matched records from the left table (`employee`). If there is no match, `NULL` is returned for the left table columns.

```
SELECT e.first_name, e.last_name, o.city, o.country
FROM employee e
RIGHT JOIN office o ON e.office_id = o.office_id;
```

Explanation:

- `SELECT e.first_name, e.last_name, o.city, o.country:`
 - Same selection of columns as before.
- `FROM employee e:`
 - The `employee` table is used as the left table (aliased as `e`).
- `RIGHT JOIN office o:`
 - Joins the `employee` table with the `office` table.
 - `RIGHT JOIN` includes all rows from the `office` table, even if there are no matching employees assigned to an office.
- `ON e.office_id = o.office_id:`
 - The join condition. If there is no matching `employee` for an `office_id`, the employee's columns (`first_name`, `last_name`) will be `NULL`.

Result:

- All offices will be listed. If an office has no employees, the employee's first name and last name will be **NULL**.

3. **IS NULL / IS NOT NULL Condition:**

- Used to filter records where certain columns are either **NULL** or not **NULL**.

Example:

Suppose you want to join all employees and return only those who have not been assigned an office (**office_id** is **NULL**).

```
SELECT e.first_name, e.last_name
FROM employee e
LEFT JOIN office o ON e.office_id = o.office_id
WHERE e.office_id IS NULL;
```

Explanation:

- **WHERE e.office_id IS NULL**: This condition filters only those employees who don't have an assigned **office_id** (i.e., the **office_id** is **NULL**).
- The **LEFT JOIN** ensures that all employees are included, and this condition filters for those with no office.

4. **FULL OUTER JOIN (Simulated with UNION)**

Returns all records when there is a match in either table. If there is no match, **NULL** is returned for the missing side. We use **UNION** to simulate a **FULL OUTER JOIN** in systems that don't support it directly

```
SELECT e.first_name, e.last_name, o.city, o.country
FROM employee e
LEFT JOIN office o ON e.office_id = o.office_id
UNION
SELECT e.first_name, e.last_name, o.city, o.country
```

```
FROM employee e
RIGHT JOIN office o ON e.office_id = o.office_id;
```

Explanation:

- **LEFT JOIN part:**
 - Retrieves all employees and their corresponding offices (if any). Employees without an office will have **NULL** for office information.
- **RIGHT JOIN part:**
 - Retrieves all offices and their corresponding employees (if any). Offices without employees will have **NULL** for employee information.
- **UNION:**
 - Combines the results from the **LEFT JOIN** and **RIGHT JOIN**. This way, all employees and all offices are included, even if they don't match.

Result:

- All employees and offices are included. If an employee has no office, the office details are **NULL**. If an office has no employees, the employee details are **NULL**.

5. CROSS JOIN

Produces the Cartesian product of the two tables. Every row from the **employee** table is combined with every row from the **office** table.

```
SELECT e.first_name, e.last_name, o.city, o.country
FROM employee e
CROSS JOIN office o;
```

Explanation:

- **CROSS JOIN** creates a combination of every employee with every office.
- There is no **ON** condition in a **CROSS JOIN**, as it simply pairs every row from **employee** with every row from **office**.

Result:

- If there are 5 employees and 3 offices, the result will have 15 rows (every employee combined with every office).

Self Join

Let's dive into the concept of a Self Join in SQL and how it works in detail. To make it clear, we will create an **employee** table and use a self join to find the manager of every employee.

Step 1: Create the **employee** Table

Here, we will define an **employee** table where each employee can have a **manager_id** (referring to another employee in the same table).

```
CREATE TABLE employee (  
    employee_id INT PRIMARY KEY, -- Unique ID for each employee  
    first_name VARCHAR(50),      -- Employee's first name  
    last_name VARCHAR(50),       -- Employee's last name  
    manager_id INT               -- Manager's ID (links to another employee_id in the same  
table)  
);
```

-- Insert records into employee table

```
INSERT INTO employee (employee_id, first_name, last_name, manager_id) VALUES  
(1, 'John', 'Doe', NULL),      -- John Doe has no manager (CEO)  
(2, 'Jane', 'Smith', 1),       -- Jane Smith reports to John Doe  
(3, 'Michael', 'Johnson', 1),  -- Michael Johnson reports to John Doe  
(4, 'Emily', 'Davis', 2),       -- Emily Davis reports to Jane Smith  
(5, 'Robert', 'Brown', 3);      -- Robert Brown reports to Michael Johnson
```

Find the Manager of Each Employee

A **self join** is when a table is joined with itself. In this case, we are joining the **employee** table with itself to match each employee with their manager.


```
SELECT e.first_name AS Employee, m.first_name AS Manager
FROM employee e
LEFT JOIN employee m ON e.manager_id = m.employee_id;
```

Explanation (Line by Line)

1. **SELECT e.first_name AS Employee, m.first_name AS Manager:**
 - We select two columns from the self-joined table:
 - **e.first_name**: Refers to the employee's first name. This is aliased as **Employee**.
 - **m.first_name**: Refers to the manager's first name. This is aliased as **Manager**.
 - The **e** and **m** aliases represent two different "instances" of the same **employee** table.
2. **FROM employee e:**
 - The first instance of the **employee** table is aliased as **e**.
 - This represents the employees whose managers we are trying to find.
3. **LEFT JOIN employee m ON e.manager_id = m.employee_id:**
 - The second instance of the **employee** table is aliased as **m**.
 - This instance represents the managers.
 - We use **LEFT JOIN** because we want to include all employees, even if they don't have a manager (like the CEO).
 - **ON e.manager_id = m.employee_id:**
 - This is the join condition.
 - For each row in the **employee** table (alias **e**), we look for a matching row in the same table (alias **m**) where **e.manager_id** matches **m.employee_id**.
 - Essentially, we are saying: "For every employee, find the row where the employee's manager ID matches the employee ID of another person."

Result (Background Mechanism)

Let's break down how SQL processes this **self join** row by row:

- **Step 1:** The **employee** table (**e**) row for **John Doe** (employee_id 1) is considered.
 - Since John has no manager (**manager_id** is **NULL**), there is no match found in the **m** table (other employees). Therefore, the result will be:
 - **Employee:** John
 - **Manager:** **NULL** (John has no manager).
- **Step 2:** The **employee** table (**e**) row for **Jane Smith** (employee_id 2) is considered.

- Jane's **manager_id** is 1, so SQL finds the row in the **employee** table (**m**) where **employee_id** is 1, which is **John Doe**.
- The result will be:
 - **Employee:** Jane
 - **Manager:** John
- **Step 3:** The **employee** table (**e**) row for **Michael Johnson** (**employee_id** 3) is considered.
 - Michael's **manager_id** is also 1, so SQL again finds the row in the **employee** table (**m**) where **employee_id** is 1, which is **John Doe**.
 - The result will be:
 - **Employee:** Michael
 - **Manager:** John
- **Step 4:** The **employee** table (**e**) row for **Emily Davis** (**employee_id** 4) is considered.
 - Emily's **manager_id** is 2, so SQL finds the row in the **employee** table (**m**) where **employee_id** is 2, which is **Jane Smith**.
 - The result will be:
 - **Employee:** Emily
 - **Manager:** Jane
- **Step 5:** The **employee** table (**e**) row for **Robert Brown** (**employee_id** 5) is considered.
 - Robert's **manager_id** is 3, so SQL finds the row in the **employee** table (**m**) where **employee_id** is 3, which is **Michael Johnson**.
 - The result will be:
 - **Employee:** Robert
 - **Manager:** Michael